

Productionizing predictive models

GoDataFest – Open Source

Niels Zeilemaker

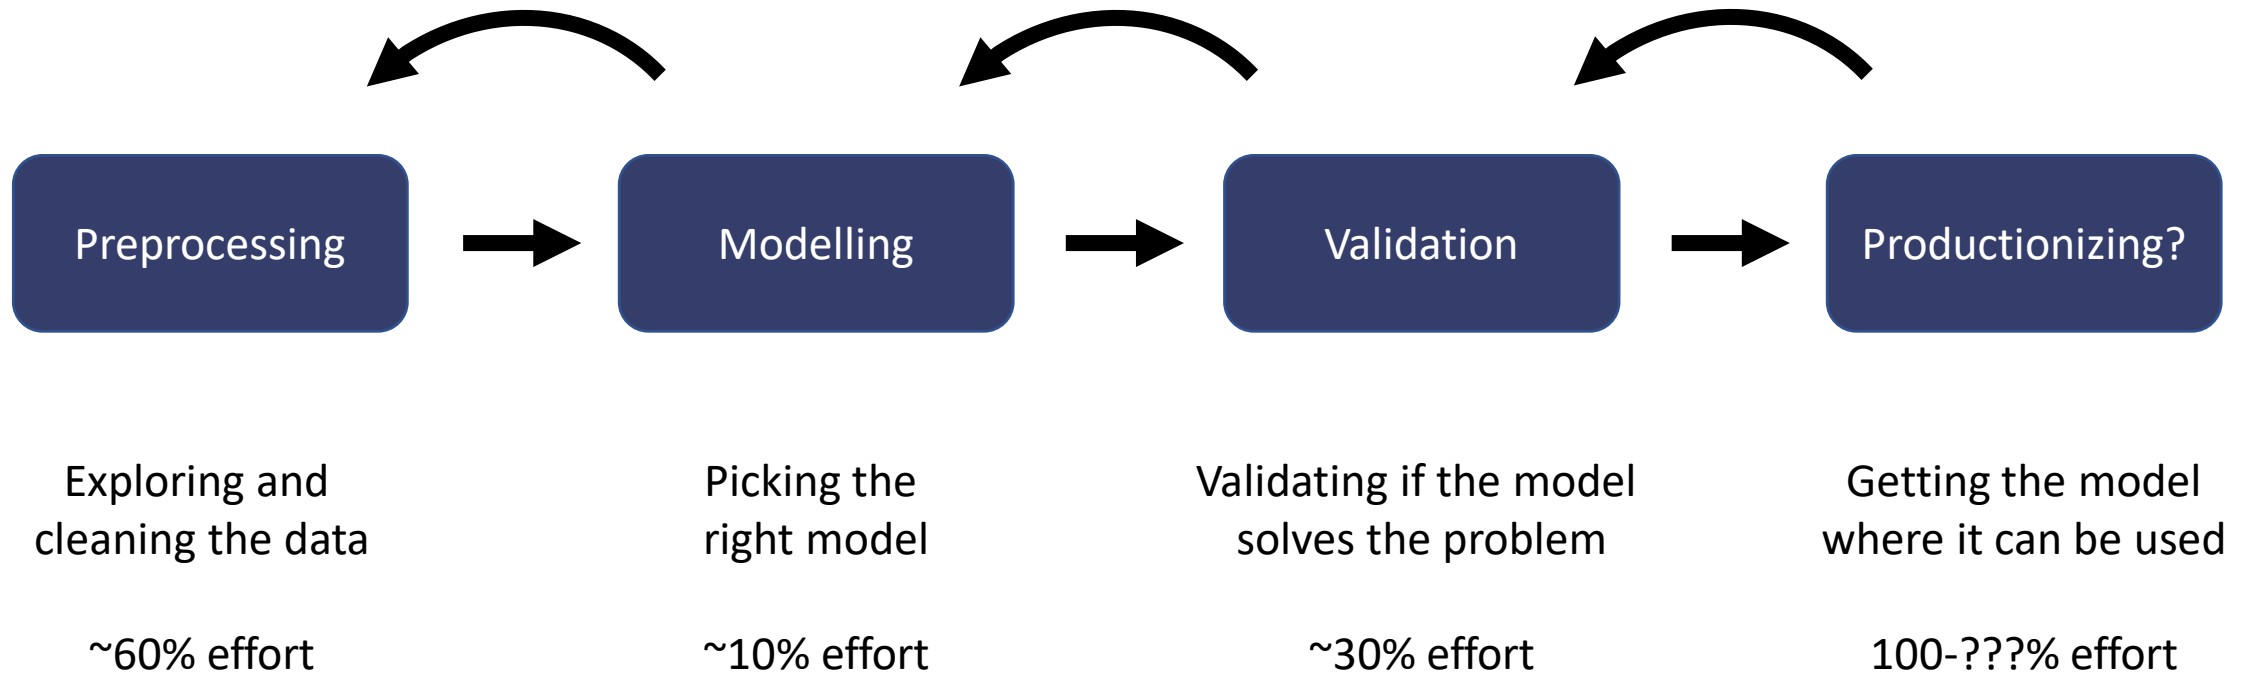
GO 
DATA
DRIVEN

This hackaton

- ~15 minutes introduction
- ~1-1,5 hours hackathon / demo
- ~15 minutes wrap-up
- Plan to finish around 11:45

The machine learning process

The machine learning process



Productionizing ML models

What is productionizing?


Productionizing

- Goal – convert model into a (standard) format that can be run in production
- How – depends on the production environment
 - Re-write code into a (production-quality) library
 - Wrap in an API for interfacing with other components

Productionizing

- Follow best practices
 - Version control
 - Code quality checks, unit testing
 - Logging / monitoring
- Consider deployment patterns
 - How will the model learn and predict?
 - What will we expose to the outside world?

Example: web-based API



However, many models look like this

Jupyter titanico-model (autosaved)

Logout

FileEditViewInsertCellKernelWidgetsHelpTrustedPython 3

SaveNewOpenRecentCopyPasteUndoRedoFindReplaceRunInterruptRestartClearOutputMarkdown

RunInterruptRestartClearOutputMarkdown

Loading the data

We can load the dataset into a dataframe by calling `pd.read_csv`. Note that we split our dataset into a train and test set, so that we can do some exploration on our training set and later gauge the accuracy of our model on the test set.

```
In [92]: 1 %matplotlib inline
          2 import pandas as pd
          3
          4 # Read train/eval datasets.
          5 data = pd.read_csv('../data/train.csv')
          6
          7 # Split data into train/test set.
          8 data_train, data_test = train_test_split(data, test_size=0.2, random_state=42)
          9
         10 data_train.sample(3, random_state=42)
```


Out[92]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
373	374	0	1	Ringhini, Mr. Sante	male	22.0	0	0	PC 17760	135.6333	NaN	C
848	849	0	2	Harper, Rev. John	male	28.0	0	1	248727	33.0000	NaN	S
593	594	0	3	Bourke, Miss. Mary	female	NaN	0	2	364848	7.7500	NaN	Q

Visualizing some variables

Visualizing data is crucial for recognizing underlying patterns to exploit in the model.

```
In [136]: 1 sns.barplot(x="Embarked", y="Survived", data=data_train);
```



How do we move this into production?

- Start building a Python package
 - Isolate main components, move these into modules
 - Identify building blocks -> make reusable functions/classes
- Improve code quality
 - Implement quality checks (pylint) and tests (pytest)
 - Document code (docstrings) and package (readme, etc.)
- Wrap model in an API (Python, Flask)

Hackathon

- Background
 - Client interested in upselling cruise ship tickets
 - Noticed that in the titanic disaster, people in higher ticket classes had a higher chance of survival
 - Would like to present this information during the booking process to sell more 1st class tickets

Hackathon

- Scenario
 - Data scientist has created a model predicted survival probabilities based on the titanic dataset
 - We have been asked to move his/her notebook into production
- Goal – build a documented + tested Python package that exposes the model as a web API

Hackathon

- Getting started
 - Go github.com/godatadriven/code-breakfast-productionizing, read the README
 - Setup a clean Python environment and install the packages in notebook/requirements.txt
 - Try running the notebook and see if you understand its contents
- Afterwards - continue with the Step 2 (see readme)

Python packaging

Python packaging

- Goal – package Python code into a redistributable package that can easily be installed in other environments
- Terminology
 - package – a directory with an `__init__.py`
 - module – a `something.py` file
 - subpackage – a package within a package

Setup.py

```
import setuptools

with open("README.md", "r") as fh:
    long_description = fh.read()

setuptools.setup(
    name="example_pkg",
    version="0.0.1",
    author="Example Author",
    author_email="author@example.com",
    description="A small example package",
    long_description=long_description,
    long_description_content_type="text/markdown",
    url="https://github.com/pypa/sampleproject",
    packages=setuptools.find_packages(),
    classifiers=[
        "Programming Language :: Python :: 3",
        "License :: OSI Approved :: MIT License",
        "Operating System :: OS Independent",
    ],
)
```

A simple example

```
├── LICENSE
├── README.md
├── my_package
│   ├── __init__.py
│   ├── my_module.py
│   ├── second_module.py
│   └── subpackage
│       ├── __init__.py
│       └── another_module.py
└── setup.py
```

```
import my_package
from my_package import my_module
from my_package import second_module
from my_package import subpackage
from my_package.sub_package import another_module
```

Additional functionality

- Besides code, Python packages also typically contain:
 - Documentation (Sphinx)
 - Unit/integration tests (Unittest library or pytest)
 - Additional readme/configuration files
- We will go into these later

Installing your package

- Your package can easily be installed using pip:
 - `pip install .`
- During development, you can use an editable install:
 - `pip install --editable .`
- This way, edits are directly reflected in your environment.
(Note: notebooks require the autoreload extension.)

Additional resources

- [Python tutorial: modules](#)
- [Python tutorial: packaging Python projects](#)
- [A tour on Python packaging](#)
- [Python's new package landscape](#)
- [Pypackage cookiecutter template](#)

Code quality

Code quality

Code is a means to communicate not only with machines but also with other developers. High quality code is good communication.

High quality code is:

- Correct runs correctly
- Human readable is easy to understand
- Consistent same formatting and naming
- Modular small units of logic
- Reusable code can be ported to/from other projects

Code style

- Every language has it's own accepted style guide(s)
 - Consistent reading experience
 - Easy to recognize what code does
 - Don't invent your own style
- Examples
 - Python – PEP8
 - R – Google Style or Advanced R

Code style – PEP8

- Examples

- Functions/variables `lower_case_variable`
- Classes `UpperCaseClass`
- Whitespace `do_this(whitespace, next, to, commas,
and, proper, indentation)`

- Many useful tools

- Style checkers – Flake8, Pylint
- Automated formatters – YAPF, Black

Programming principles

- Do one thing and do it well
 - Have small, focused functions/classes that only do one thing
 - Functions should be logical units



```
def detect_machines(data, start, end):  
    # Filter data  
    dates = pd.date_range(start, end)  
    start = start - 1  
    end = end + 1  
    filtered = data[start:end]  
    # Find events  
    ...  
    # Count machines  
    ...
```



```
def detect_machines(data, start, end):  
    filtered = filter_data(data[start:end])  
    machines = find_machines(filtered)  
    n_machines = count_machines(machines)  
    return n_machines  
  
def filter_data(start, end):  
    ...  
  
def detect_events(filtered):  
    ...  
  
def count_machines(machines):  
    ...
```

Programming principles

Goal – write code other people (and future you) can understand

Don'ts

- Long functions doing multiple things
- Copy/paste code
- Re-use variable names in function
- One-char variables, abbreviations
- Variables that differ by one character
- Long, complicated variable names
- Many temp vars (or using your own defaults)
 - `bassie, buh, zip`

Do's

- Small functions doing one thing
 - `check_boiler()`
 - `load_rankings()`
- Build libraries, functions, classes
- Follow existing design patterns
- Descriptive and concise variables:
 - `male_user, is_fridge`
- Common temp vars
 - `temp, df`

Additional resources

Software development skills for data scientists:

- http://treycausey.com/software_dev_skills.html

Machine learning in production

- <http://www.slideshare.net/turi-inc/machine-learning-in-production>

Some Design Patterns for Real World Machine Learning Systems

- <http://www.slideshare.net/justinbasilico/is-that-a-time-machine-some-design-patterns-for-real-world-machine-learning-systems>

Documentation

GO 
DATA
DRIVEN

Documentation

- One way to improve the readability of your code is to add (proper) documentation
- Different documentation types
 - Inline comments - explain what specific pieces of code do
 - Docstrings - document Python functions/classes/modules
 - Actual documentation – how to install, usage guide, etc.

Comments

- Avoid adding comments explaining the obvious
- Think about the choices/assumptions you make in your code, which are not directly clear from the code itself



```
# import packages  
import pandas as pd  
  
# load some data  
df = pd.read_csv('data.csv', skiprows=2)
```



```
# Data contains two lines of description  
# text, skip to avoid errors.  
df = pd.read_csv('data.csv', skiprows=2)
```

Docstrings

- Docstrings document how to use specific functionality

```
def rescale_between(array1d, lower=0.0, upper=5.0):  
    """Rescales array values between given upper/lower bounds.  
  
    :param np.ndarray array1d: Values to be rescaled.  
    :param float lower: Lower bound of the rescaled values.  
    :param float upper: Upper bound of the rescaled values.  
  
    :returns: Array containing rescaled values.  
    :rtype: np.ndarray  
    """  
    ...
```

- Can be accessed using `help(...)` (or ? in IPython/Jupyter)

Sphinx



- Sphinx is the de-facto tool to use in Python for writing and generating docs
- Docs are written in [reStructuredText](#)
- Many features
 - Hierarchical structure, table of contents, etc.
 - Generating docs from docstrings
 - Different themes, output types (html)

Sphinx – getting started



- [Getting started](#)
 - Install sphinx using ``pip install sphinx``
 - Generate initial template using ``sphinx-quickstart``
 - Start editing your docs
 - Build your docs using ``make html`` (in the docs folder)
- Some templates (cookiecutter-pypackage) include an initial structure that you can use

Testing

Testing

- Why test?
 - Confirm that your code does what you expect
 - Prevent regressions (code changes that change behavior)
- Two (main) types of tests
 - Unit tests – tests a single function/method
 - Integration tests – tests behavior of combined functions

Testing frameworks

- Python has multiple testing frameworks
 - unittest – builtin framework, inspired by JUnit
 - Nose/Pytest – popular third party libraries
- We will focus on [Pytest](#)
 - Easy to use, with little boilerplate code
 - Can be a bit ‘magic’ in the beginning

Test structure

```
├── my_package
|   ├── __init__.py
|   ├── helpers.py
|   └── utils.py
└── tests
    ├── conftest.py
    └── my_package
        └── test_helpers.py
├── setup.py
└── ...
```

Package code

Tests mirror
package structure

A simple example

```
# test_helpers.py

from my_package.helpers import add_two

class TestAddTwo:
    """Tests for the add_two helper function."""

    def test_positive(self):
        """Tests addition with a positive number."""
        assert add_two(1) == 3

    def test_negative(self):
        """Tests addition with a negative number."""
        assert add_two(-3) == -1
```

Fixtures

- Fixtures allow you to define functions that setup elements required by (multiple) tests

```
import pytest
import smtplib

@pytest.fixture(scope="module")
def smtp_connection():
    return smtplib.SMTP("smtp.gmail.com", 587, timeout=5)

def test_smtp(smtp_connection):
    ...
```


Running pytest

```
$ pytest
```

```
===== test session starts =====
```

```
platform linux -- Python 3.x.y, pytest-3.x.y, py-1.x.y, pluggy-0.x.y
```

```
rootdir: $REGENDOC_TMPDIR, inifile:
```

```
collected 1 item
```

```
test_sample.py F [100%]
```

```
===== FAILURES =====
```

```
_____ test_answer _____
```

```
def test_answer():
```

```
> assert inc(3) == 5
```

```
E assert 4 == 5
```

```
E + where 4 = inc(3)
```

```
test_sample.py:6: AssertionError
```

```
===== 1 failed in 0.12 seconds =====
```

When to stop testing?

- So when do we have enough tests?
 - Ideally – when our code is bug-free
 - In practice – when we have ‘enough’ confidence in our code
- A popular metric is code coverage
 - Percentage of code covered by tests
 - Note: code with 100% coverage is not bug-free
- Can be generated in pytest using `pytest-cov` plugin

Building web API's using Flask

Web theory – methods

- There are different types of request methods:
 - GET - Retrieve the resource from the server
 - POST - Create a resource on the server
 - PUT - Update the resource on the server
 - DELETE - Delete the resource from the server
- In general; you should keep GET requests limited to requests that do not change the state of the server.

Web theory – response types

- The status of a HTTP request is indicated using a code:
 - 1xx - continue
 - 2xx - you got a response
 - 4xx - server thinks a client made an error
 - 3xx - redirect
 - 5xx - server thinks that it made an error
- We're omitting a lot of details now and a full summary can be found [here](#).

Flask

“Flask is a microframework for Python based on Werkzeug, Jinja 2 and good intentions.”



- A minimal [example](#):

```
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello_world():
    return 'Hello, World!', 200
```

```
$ export FLASK_APP=hello.py
$ python -m flask run
```

Flask – class-based approach

```
from flask import Flask

class App(Flask):
    def __init__(self, *args, **kwargs):
        super().__init__(*args, **kwargs)
        self.add_url_rule(
            "/",
            view_func=self.hello_world,
            methods=["GET"])

    def hello_world(self):
        return "Hello world!", 200
```

Flask is single-threaded

- Flask starts a Python process which is a single thread.
This means it can only handle one request at a time.
- You must wrap the application in a WSGI (Web Server Gateway Interface) to serve multiple clients simultaneously.
- For more details, take a look at [deploying Flask in Production](#).

Making requests

“Requests is the only Non-GMO HTTP library for Python, safe for human consumption.”



```
>>> r = requests.get('https://api.github.com/user')
>>> r.status_code
200
>>> r.headers['content-type']
'application/json; charset=utf-8'
>>> r.text
'{"type": "User" ... '
>>> r.json()
{'private_gists': 419, 'total_private_repos': 77, ...}
```

GO 
DATA
DRIVEN

Containerization using Docker

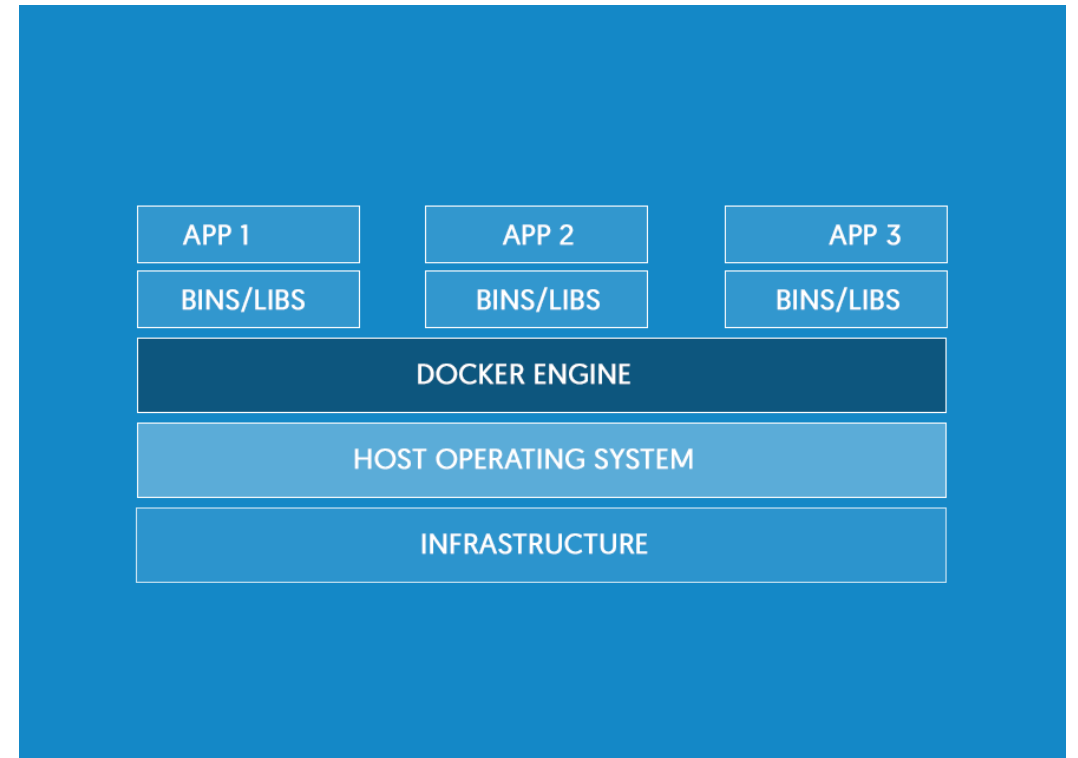
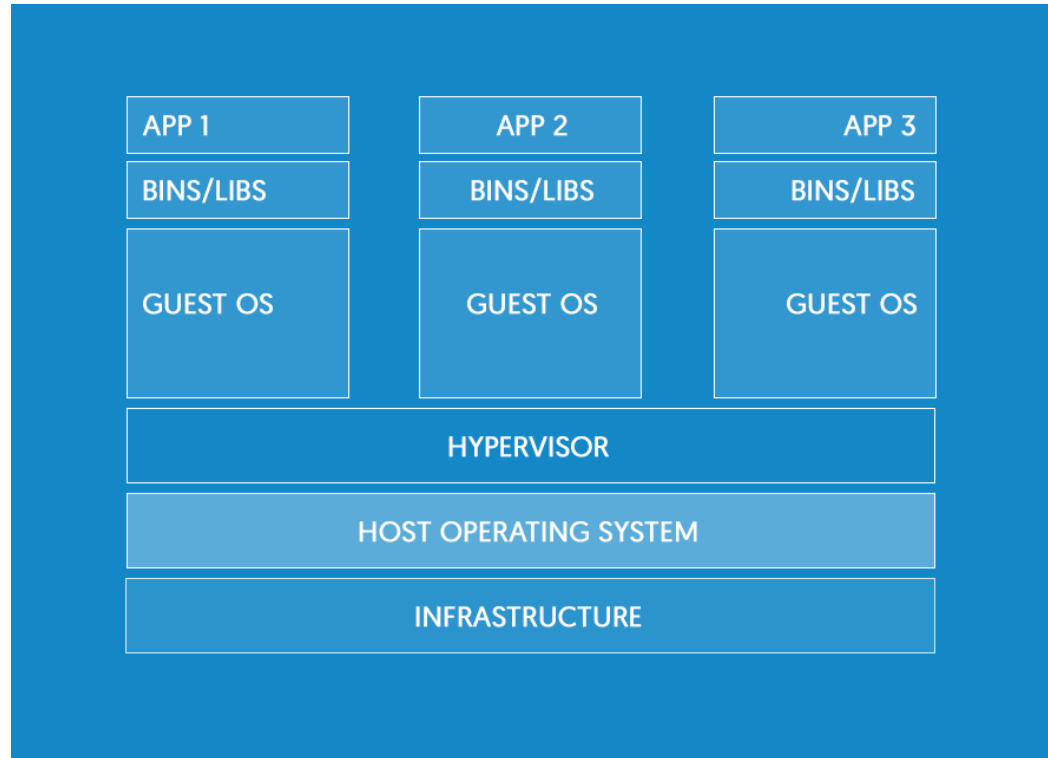
Why containers?

- Package your application
- Run it everywhere
- Lightweight
- No more dependency hell


Containers vs. VM's

- VMs run on top of a hypervisor
 - Includes an OS
 - Dependencies
 - Your application
- Containers
 - Run on a layered filesystem
 - Run as an isolated process
 - Share the kernel

Containers vs. VM's



Docker



- Most popular container engine
- Provides an easy CLI to create and manage containers
- Rapid development
 - New version every ~3 months

Dockerfile example

```
FROM ubuntu:16.04
MAINTAINER Sven Dowideit <SvenDowideit@docker.com>

RUN apt-get update && apt-get install -y openssh-server
RUN mkdir /var/run/sshd
RUN echo 'root:screencast' | chpasswd
RUN sed -i 's/PermitRootLogin prohibit-password/PermitRootLogin yes/' /etc/ssh/sshd_config

# SSH login fix. Otherwise user is kicked off after login
RUN sed 's@session\s*required\s*pam_loginuid.so@session optional pam_loginuid.so@g' -i /etc/pam.d/sshd

ENV NOTVISIBLE "in users profile"
RUN echo "export VISIBLE=now" >> /etc/profile

EXPOSE 22
CMD ["/usr/sbin/sshd", "-D"]
```

Questions?